

---

# Beyond containers:

The value of clusters and microservices

---

Craig McLuckie

Group Product Manager, Google Corporation

*Founder, Kubernetes Project*

*Chairman, Cloud Native Computing Foundation*




# History of Containers

- Control Groups (cgroups)
  - Support efficient resource isolation for process groups
  - Contributed to Linux by Google (2006)
- Namespaces
  - Isolates global system resources
  - Supports OS level virtualization
- LXC
  - Lightweight userland interface to Linux kernel container features
  - Offers isolation and namespace support using cgroups
  - Barebones, but stable toolset
  - Early use was in PaaS platforms
    - efficient multi-tenant scheduling (somewhat weak security)



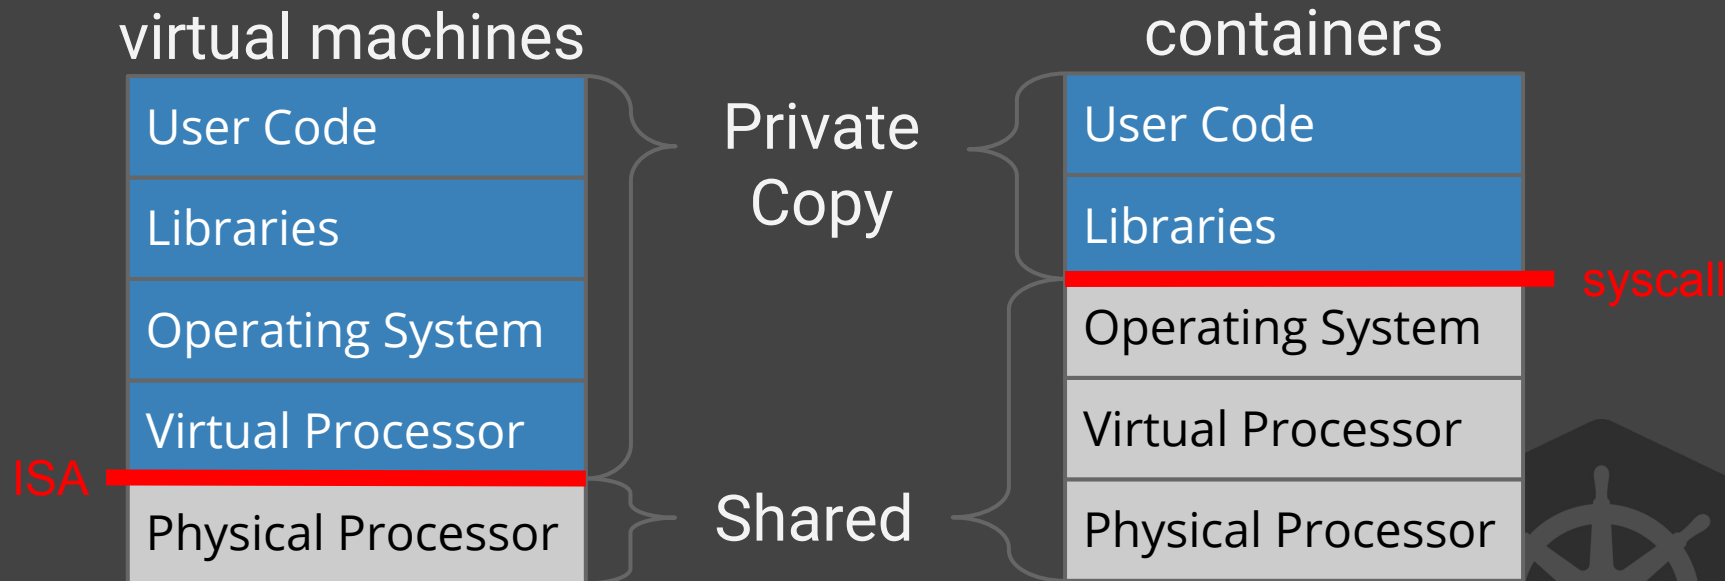
# Docker



- Docker wraps everything your code needs to run in a complete filesystem.
- Docker is magic in three ways:
  - Real portability: Syscall layer is same everywhere
  - High degree of reuse and extensibility: Stackable file system
  - Easy to use: Simple and accessible tooling



# Container vs VM



**containers: less overhead, enable more “magic”**

# So what remains?

---

- Development is a big part of the story, but what about operations?
  - Can we make better use of resources?
  - Can we reduce the cost of operations?
  - Can we build more flexible systems?
- Can we make it easier to build and run distributed systems?



## Deploy to cluster not machine

- A cluster is a collection of machines managed as a single unit
- In a cluster...
  - a microservice is an atom of software management and consumption
  - deployment is managed for you
  - your application cannot be affected by the health of a single machine




# Kubernetes

- An Open Source technology started by Google
- Links together a number of virtual or physical machines to create a cluster
- Fills the operations gap for containers
  - Deploy container
  - Scale and manage health of containers
  - Connect containers to network
  - Attach storage
- Creates programmable 'logical infrastructure'



# Kubernetes pod


- Unit of deployment
- One or more containers
- Shared fate
- Mapped dynamically to a host node
- Assigned a network interface
- Assigned volume(s)





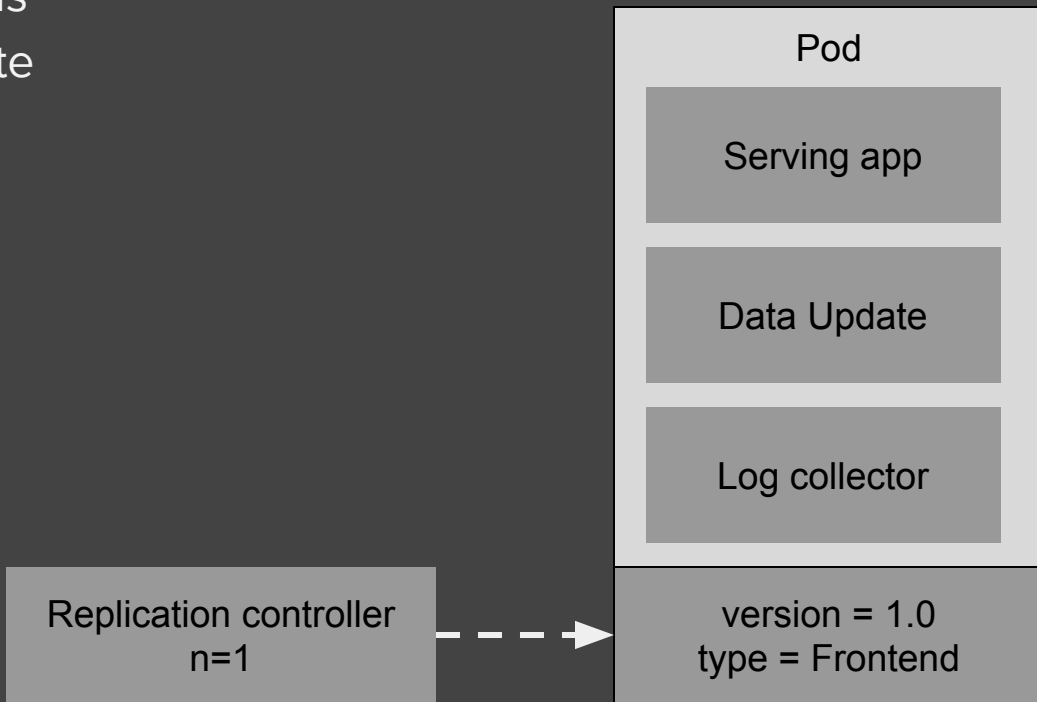
# Kubernetes labels

- Label anything
- Name-value pair
- Create your own
- Membership determined via a label selector



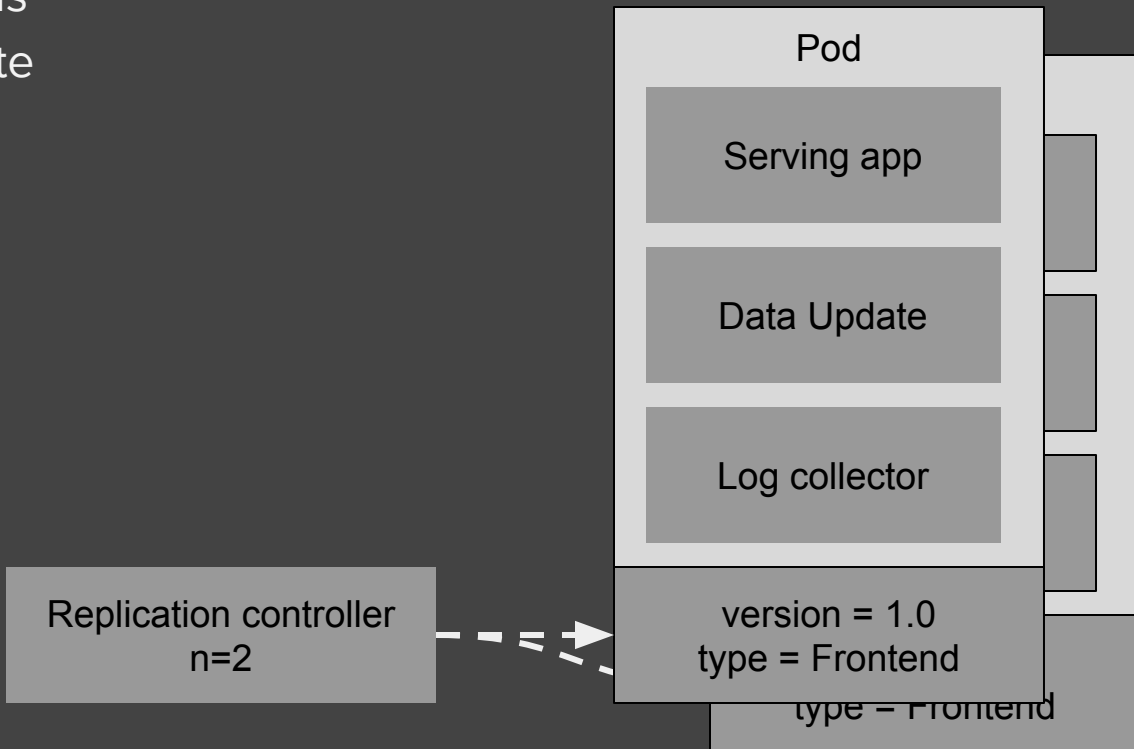
# Replication controller

- Manage the lifecycle of pods
- Drive system to desired state
- If too few add pods
- If too many kill pods



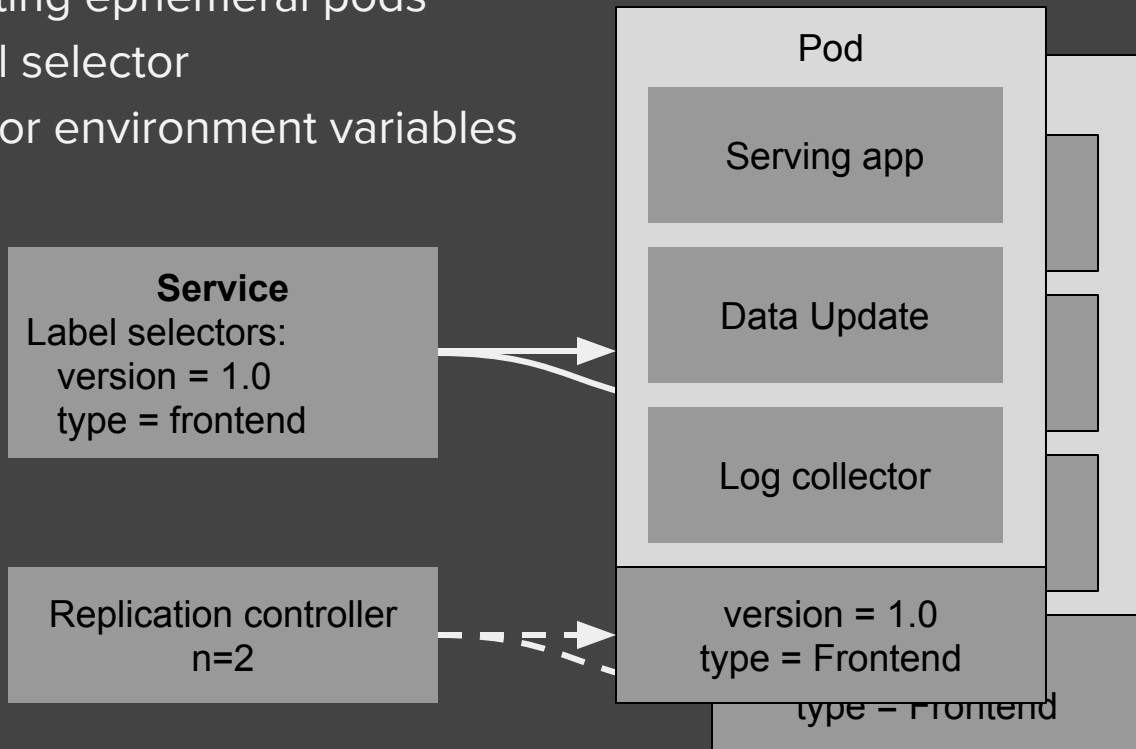
# Replication controller

- Manage the lifecycle of pods
- Drive system to desired state
- If too few add pods
- If too many kill pods




# Kubernetes service

- Durable endpoint representing ephemeral pods
- Membership driven by label selector
- Discoverable through DNS or environment variables



# Kubernetes namespaces

- Allow teams/services to be isolated
- Enable reuse of configuration
- Logical partition for quota and authZ



# Kubernetes services and operations

- Label selector gives programmatic control over membership
- Service is a minimum atom of software consumption in a cluster
  - Minimally (hostname, port)
  - Optionally services can consume Endpoint API to understand micro-service membership
- Natural path to richer orchestration
  - Route based on label selector
  - Blue/Green
  - Canary
- Service offers scoping options
  - Cluster level common services
  - Namespace scope services
- 'Mix in' external services by creating a Service object without a selector
  - 'Native' consumption experience



# SysAdmins

## Developers

- Code

- + More predictable
- + A first layer of specialization
- + Auditable, programmatic processes

## Sys admin and release managers

- Install and configure

- Slow
- Scales sub-linearly
- Human operator

Granularity of operation: **Ticket**



# DevOps

Developers can code, right?

- Business logic in language A
- Deployment logic in language B

Granularity of operation: **Integration**

- + Hardware ops goes away
- + Imperative repeatable process
- + Faster deployments
- Doubling the necessary skillset
- Running imperative code in prod
- Effort scales ~linearly
- Loss of global control/visibility






# Cloud Native Ops

Systems rely common services and specialized teams

- Describe and package app
- Programmatic infrastructure
- Policy based on logical units
- Predictable common services

Granularity of operation: **API call**

- + Scales efficiently
- + Tools align with teams
- + Perfectly predictable and repeatable
  
- Requires cultural change
- Nascent: Still missing some tools
- Nascent: Missing stable interfaces



## Common services

- Application services
  - Implementation detail (breaking down the monolith)
  - Shared artifacts, private instance
  - Multi-tenant
    - Private/Soft (you trust/control users)
    - Public/Hard (your users are hostile/idiots)
- Common distributed system services
  - Naming, discovery, quoruming, data sharding, etc generally just part of your cluster



## Specialized roles emerge

- Infrastructure Operator
  - Generally your cloud provider
- Cluster and Services Operators
  - **Engineers** (SRE) run clustering technology, and support common cluster services
- Application operator
  - **Engineers** (SRE) run applications
- Application developer

This sounds like devops, and it could be. Difference is the tools are better, and there is a clear path to specialization. SREs automate themselves out of a job.

These are “hats” that people wear, not job titles.



# An obvious tension

- Running things tend to keep running (mostly)
- Change breaks things
  - Binaries
  - Dependencies
  - Flags
- But the business is never going to sit still
  - (Software is eating the world)
- **Managing down the cost of change** is a critical competitive factor
  - Reduce toil
  - Increase control
  - Improve recoverability
  - Improve velocity



# Architecting for change

- Adopting modern CI/CD practices solve a local problem
- Putting shared things behind a stable interface
- Track, tool and automate change
  - Declarative trumps imperative
  - Dead simple updates and rollback
- Remove degrees of freedom as early as possible in the deployment cycle
  - Static binary/image, fully resolved configuration, etc
  - Steer clear of imperative steps in production
  - Minimize deltas between dev, staging, prod
- Run multiple versions of the same thing in production
  - Learn small
  - Shift load carefully
  - Run experiments



## Scaling your engineering team

- Remember Conway's law
- Microservices allow a 'crew' to solve a specific problem
- 'Throwing product over the wall' → 'running a service'
- Common services allow focus on just business logic



# Driving up operational maturity

Specialization of operations can create truly differentiated operations professionals

1. Service level monitoring
  - a. Latency
  - b. Traffic
  - c. Errors
  - d. Saturation
2. Incident response playbooks
3. Blameless post-mortem / RCA
4. Qualification and staged releases
5. Experiments



# Community and Cloud Native Computing Foundation

- Kubernetes is not just a Google project
- Kubernetes doesn't live alone -- it is an important piece of the bigger story
- We need a safe place to collaborate and innovate
- That is why we donated Kubernetes to Cloud Native Computing Foundation



**CLOUD NATIVE**  
**COMPUTING FOUNDATION**

