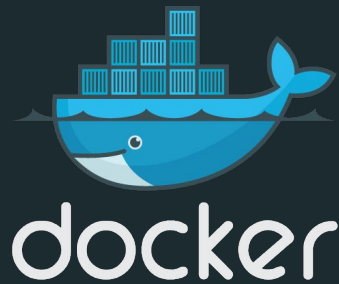


# What's New in Docker 1.12

(Spoiler alert: a lot!)

Mike Goelzer



# Swarm Mode





Engine



```
$ docker swarm init
```


# Swarm Mode




 `$ docker swarm init`

 `$ docker swarm join <IP of manager>:2377`


# Swarm Mode



 `$ docker swarm init`


 `$ docker swarm join <IP of manager>:2377`

# Services



```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp frontend_image:latest
```

# Services





```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp frontend_image:latest
```




```
$ docker service create --name redis --network mynet redis:latest
```


# Node Failure





```
 $ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp frontend_image:latest
```

```
 $ docker service create --name redis --network mynet redis:latest
```

# Node Failure




 `$ docker service create --replicas 3 --name frontend --network mynet --publish 80:80/tcp frontend_image:latest`

 `$ docker service create --name redis --network mynet redis:latest`



# Desired State $\neq$ Actual State




```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp frontend_image:latest
```



```
$ docker service create --name redis --network mynet redis:latest
```

# Converge Back to Desired State

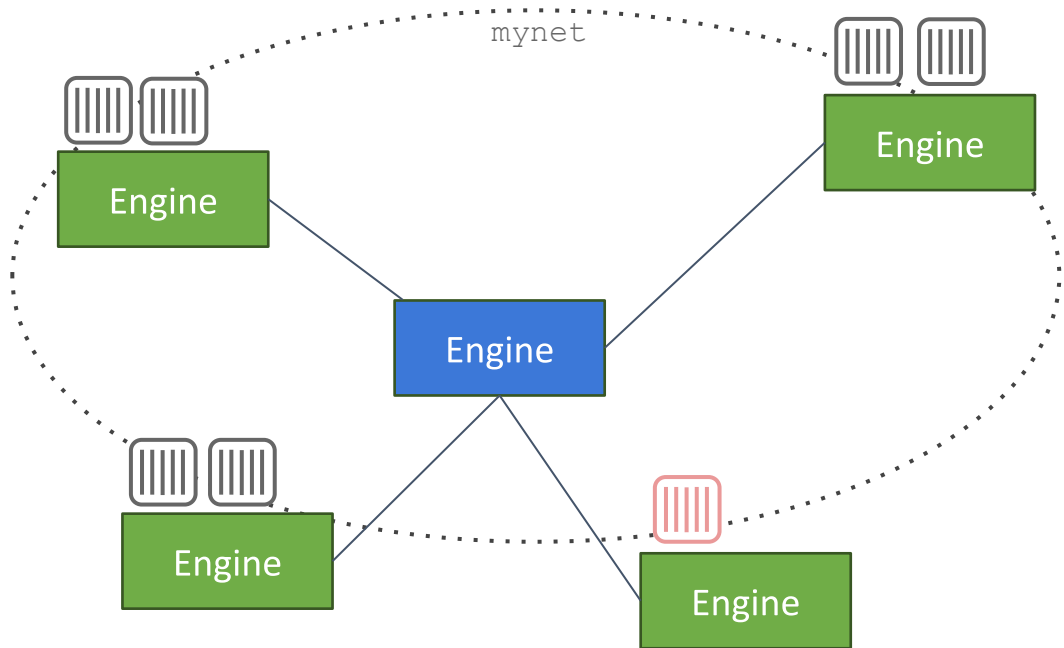


```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp frontend_image:latest
```




```
$ docker service create --name redis --network mynet redis:latest
```

# Scaling




```
$ docker service scale frontend=6
```


# Scaling




```
$ docker service scale frontend=10
```

# Global Services




 `$ docker service create --mode=global --name prometheus prom/prometheus`

# Constraints




# Constraints



```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp --constraint com.example.storage="ssd"  
frontend_image:latest
```


# Constraints





```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp --constraint com.example.storage="ssd"  
frontend_image:latest  
$ docker service scale frontend=10
```




# Services



# Services are grouped into stacks



# Distributed Application Bundle (.dab) declares a stack

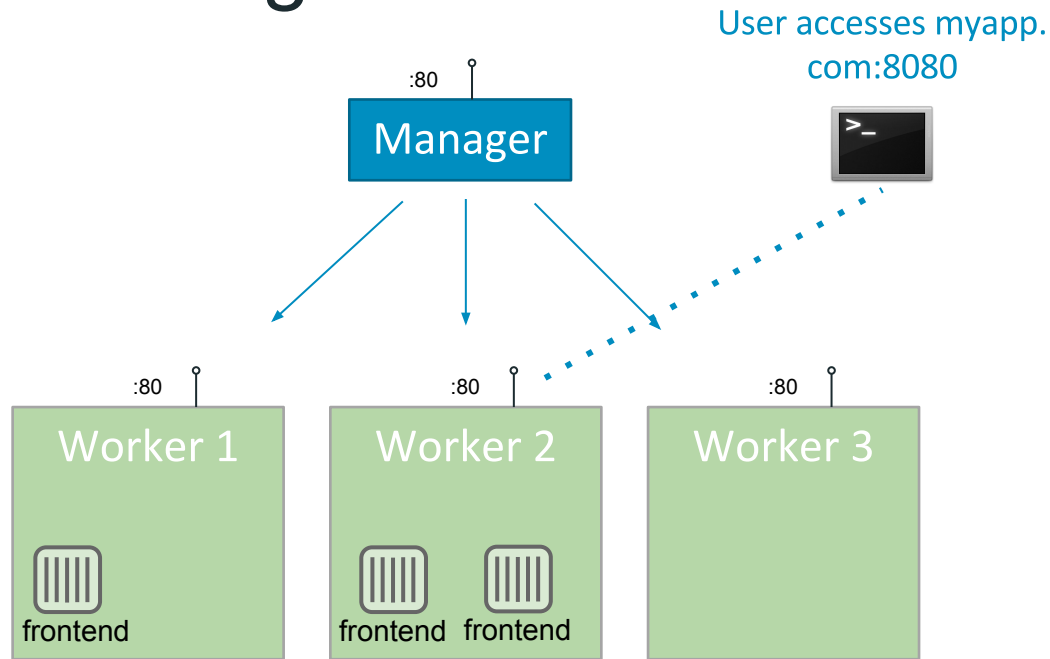


# Swarm mode orchestration is optional

- You don't have to use it
- 1.12 is fully backwards compatible
- Will not break existing deployments and scripts



# Routing Mesh



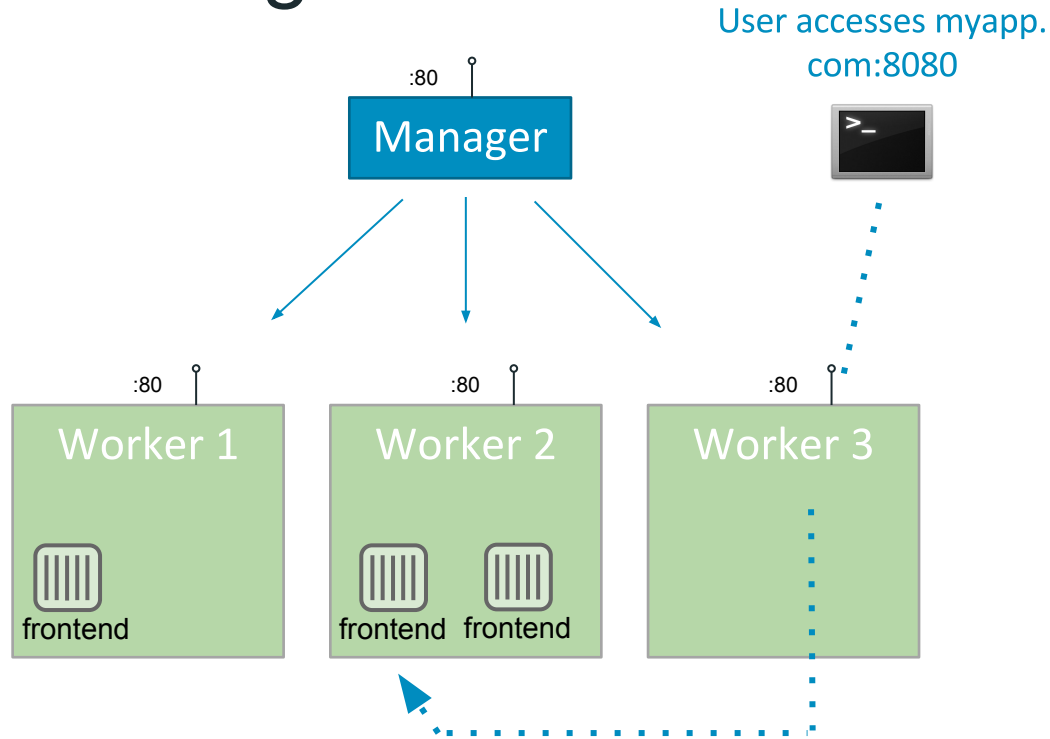
- Operator reserves a swarm-wide ingress port (80) for myapp
- Every node listens on 80
- Container-aware routing mesh can transparently reroute traffic from Worker3 to a node that is running container
- Built in load balancing into the Engine
- DNS-based service discovery



```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp frontend_image:latest
```



# Routing Mesh: Published Ports



- Operator reserves a swarm-wide ingress port (80) for myapp
- Every node listens on 80
- Container-aware routing mesh can transparently reroute traffic from Worker3 to a node that is running container
- Built in load balancing into the Engine
- DNS-based service discovery



```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp frontend_image:latest
```



# Security out of the box

- **Cryptographic Node Identity**
  - Workload segregation (think PCI)
- There is no “insecure mode”:
  - TLS mutual auth
  - TLS encryption
  - Certificate rotation



# Container Health Check in Dockerfile

```
HEALTHCHECK --interval=5m --timeout=3s  
  --retries 3  
  CMD curl -f http://localhost/ || exit 1
```

Checks every 5 minutes that web server can return index page within 3 seconds.

Three consecutive failures puts container in an unhealthy state.





# New Plugin Subcommands

```
docker plugin install tiborvass/no-remove
```

```
docker plugin enable no-remove
```

```
docker plugin disable no-remove
```



# Plugin Permissions Model

```
$ docker plugin install tiborvass/no-remove
```

```
Plugin "mikegoelzer/myplugin:latest"  
requested the following privileges:
```

- Networking: host
- Mounting host path: /data


```
Do you grant the above permissions? [y/N]
```




# Orchestration Deep Dive

DockerCon 2016


Andrea Luzzardi



# Swarm Topology



# Swarm Topology




Manager

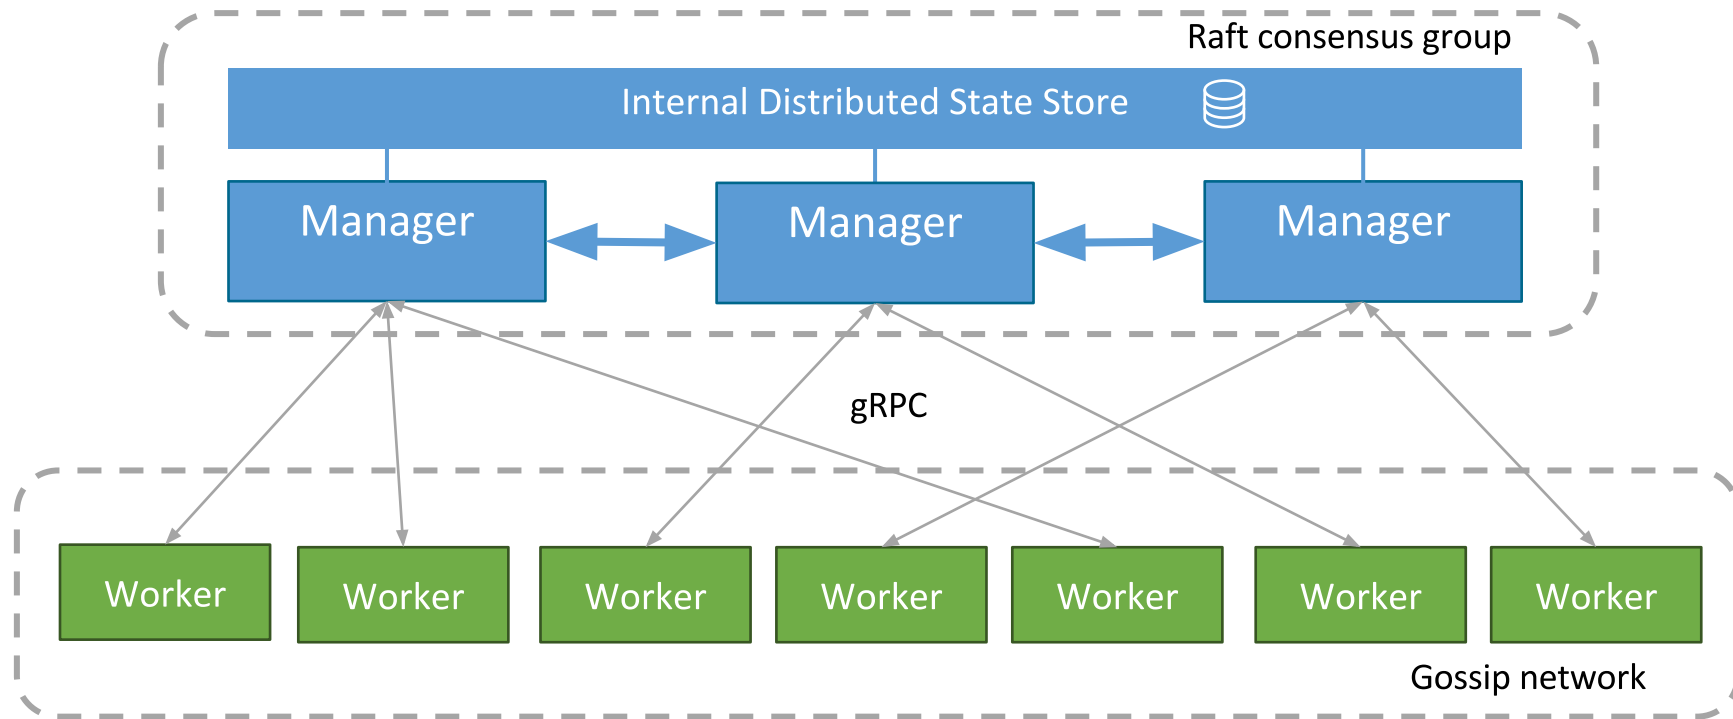
Worker

# Swarm Topology

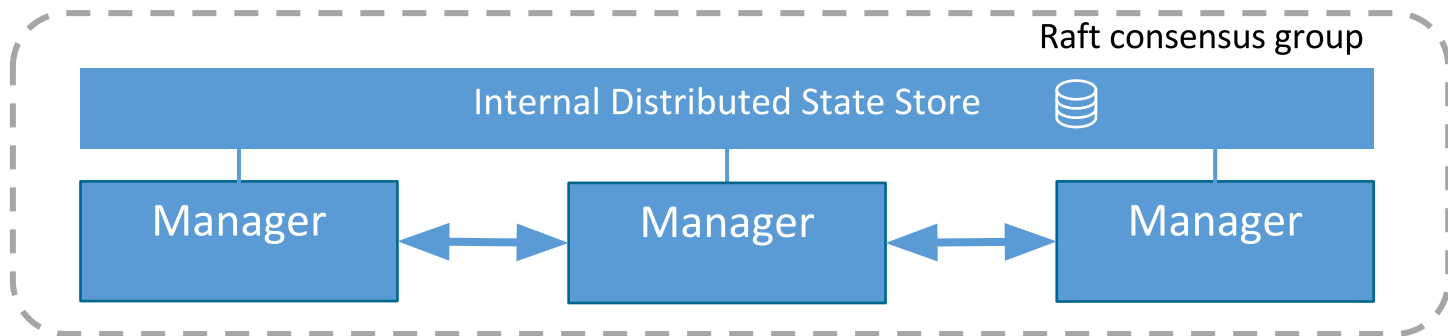
- Each Node has a role
- Roles are dynamic
- Programmable Topology



# Docker Swarm Communication Internals



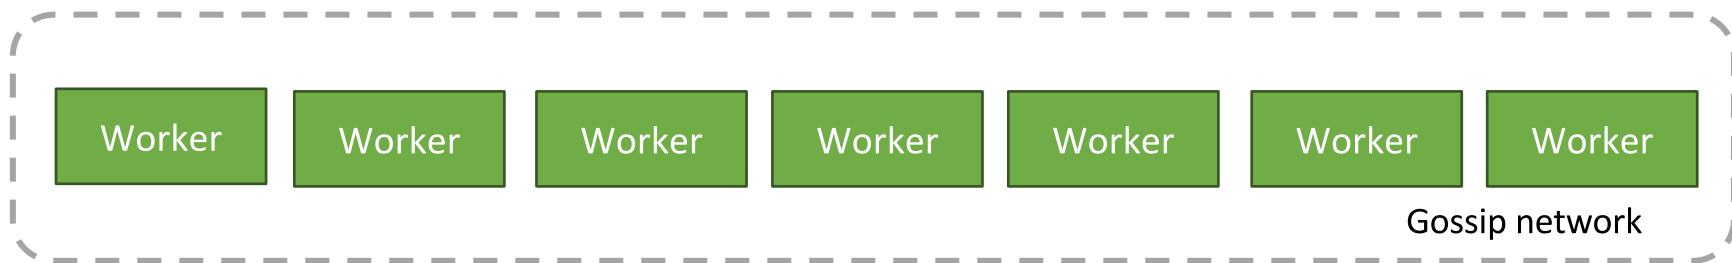
# Quorum Layer



- Strongly consistent: Holds desired state
- Simple to operate
- Blazing fast (in-memory reads, domain specific indexing, ...)
- Secure

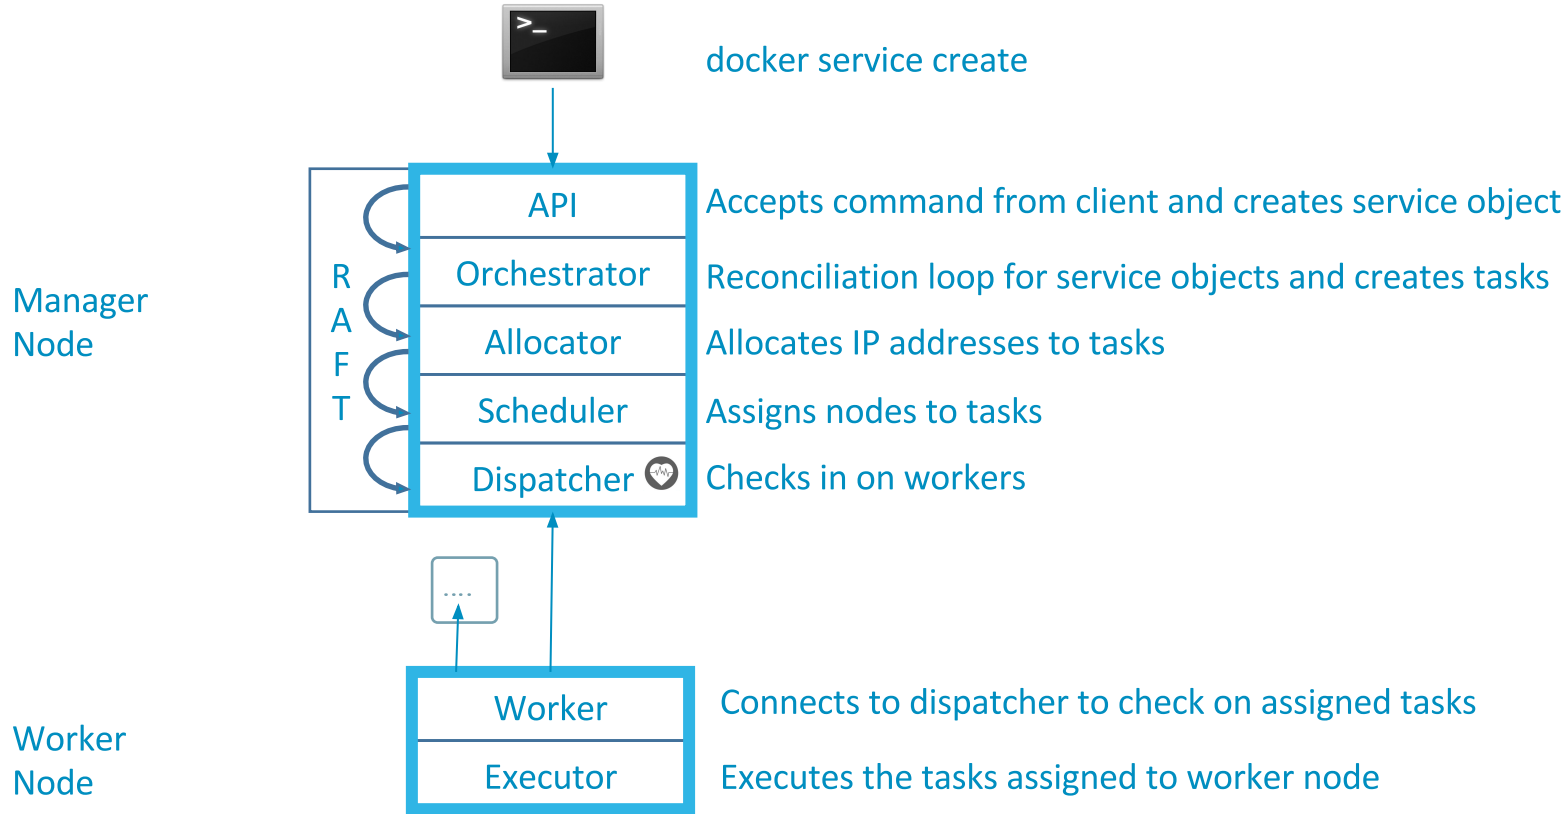


# Worker-to-Worker Gossip




- Eventually consistent: Routing mesh, load balancing rules, ...
- High volume, p2p network between workers
- Secure: Symmetric encryption with key rotation in Raft

# Node Breakdown




# Internal Load-Balancer




- **Load-balancer is designed as an integral part of CNM**
  - Works on top of CNM constructs (network, endpoint, sandbox, SD)
  - Every Service gets a Virtual-IP
  - Built-in SD resolves Service-Name -> VIP
  - Service VIP -> Container IP load balancing achieved using IPVS

# Ingress Load-Balancer



# Secure by default with end to end encryption



- Cryptographic node identity
- Automatic encryption and mutual auth (TLS)
- Automatic cert rotation
- External CA integration



# Learn more about 1.12

## **Monday 5:20 pm @ Ballroom 6E**

- Docker Security Deep Dive

## **Tuesday 3:55 pm @ Ballroom 6E**

- Docker for Ops: Networking Deep Dive, Considerations and Troubleshooting




# Questions?

Mike Goelzer

[mgoelzer@docker.com](mailto:mgoelzer@docker.com) / @mgoelzer

Andrea Luzzardi

[al@docker.com](mailto:al@docker.com) / @aluzzardi



# Thank You

Mike Goelzer

[mgoelzer@docker.com](mailto:mgoelzer@docker.com) / @mgoelzer

Andrea Luzzardi

[al@docker.com](mailto:al@docker.com) / @aluzzardi

